

Problem Statement

Predict whether a driver will be leaving the company or not based on their attributes like

- Demographics (city, age, gender etc.)
- Tenure information (joining date, Last Date)
- Historical data regarding the performance of the driver (Quarterly rating, Monthly business acquired, grade, Income)

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # metrics
from sklearn.metrics import recall_score
from sklearn.metrics import f1_score
from sklearn.metrics import roc_curve
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
```

```
In [3]: df=pd.read_csv("")
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Unnamed: 0	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	Dateofjoining	LastWorkingDate	Joining Designation	Grade	Total Business Value
0	0	01/01/19	1	28.0	0.0	C23	2	57387	24/12/18	NaN	1	1	2381060
1	1	02/01/19	1	28.0	0.0	C23	2	57387	24/12/18	NaN	1	1	-665480
2	2	03/01/19	1	28.0	0.0	C23	2	57387	24/12/18	03/11/19	1	1	0
3	3	11/01/20	2	31.0	0.0	C7	2	67016	11/06/20	NaN	2	2	0
4	4	12/01/20	2	31.0	0.0	C7	2	67016	11/06/20	NaN	2	2	0

```
In [5]: df.shape
```

```
Out[5]: (19104, 14)
```

```
In [6]: df.columns
```

```
Out[6]: Index(['Unnamed: 0', 'MMM-YY', 'Driver_ID', 'Age', 'Gender', 'City',  
              'Education_Level', 'Income', 'Dateofjoining', 'LastWorkingDate',  
              'Joining Designation', 'Grade', 'Total Business Value',  
              'Quarterly Rating'],  
              dtype='object')
```

In [7]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19104 entries, 0 to 19103
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            19104 non-null  int64
1   MMM-YY               19104 non-null  object
2   Driver_ID            19104 non-null  int64
3   Age                  19043 non-null  float64
4   Gender               19052 non-null  float64
5   City                 19104 non-null  object
6   Education_Level      19104 non-null  int64
7   Income               19104 non-null  int64
8   Dateofjoining        19104 non-null  object
9   LastWorkingDate      1616 non-null   object
10  Joining Designation  19104 non-null  int64
11  Grade                19104 non-null  int64
12  Total Business Value 19104 non-null  int64
13  Quarterly Rating     19104 non-null  int64
dtypes: float64(2), int64(8), object(4)
memory usage: 2.0+ MB
```

```
In [8]: df.describe()
```

```
Out[8]:
```

	Unnamed: 0	Driver_ID	Age	Gender	Education_Level	Income	Joining Designation	Grade	Total Business Value
count	19104.000000	19104.000000	19043.000000	19052.000000	19104.000000	19104.000000	19104.000000	19104.000000	1.910400e+04
mean	9551.500000	1415.591133	34.668435	0.418749	1.021671	65652.025126	1.690536	2.252670	5.716621e+05
std	5514.994107	810.705321	6.257912	0.493367	0.800167	30914.515344	0.836984	1.026512	1.128312e+06
min	0.000000	1.000000	21.000000	0.000000	0.000000	10747.000000	1.000000	1.000000	-6.000000e+06
25%	4775.750000	710.000000	30.000000	0.000000	0.000000	42383.000000	1.000000	1.000000	0.000000e+00
50%	9551.500000	1417.000000	34.000000	0.000000	1.000000	60087.000000	1.000000	2.000000	2.500000e+05
75%	14327.250000	2137.000000	39.000000	1.000000	2.000000	83969.000000	2.000000	3.000000	6.997000e+05
max	19103.000000	2788.000000	58.000000	1.000000	2.000000	188418.000000	5.000000	5.000000	3.374772e+07

```
In [9]: lis=['City','Grade','Joining Designation','Education_Level','Gender','Quarterly Rating']
```

```
for i in lis:
    print(i, ': ',df[i].unique())
    print()
```

```
City : ['C23' 'C7' 'C13' 'C9' 'C11' 'C2' 'C19' 'C26' 'C20' 'C17' 'C29' 'C10'
        'C24' 'C14' 'C6' 'C28' 'C5' 'C18' 'C27' 'C15' 'C8' 'C25' 'C21' 'C1' 'C4'
        'C3' 'C16' 'C22' 'C12']
```

```
Grade : [1 2 3 4 5]
```

```
Joining Designation : [1 2 3 4 5]
```

```
Education_Level : [2 0 1]
```

```
Gender : [ 0.  1. nan]
```

```
Quarterly Rating : [2 1 4 3]
```

Data Preprocessing-1

Dropping unwanted column

```
In [10]: df.drop(columns='Unnamed: 0', inplace=True)
```

KNN Imputing Data-Handling Missing Value

```
In [11]: df.isnull().sum()
```

```
Out[11]: MMM-YY                0  
Driver_ID                    0  
Age                          61  
Gender                       52  
City                         0  
Education_Level              0  
Income                       0  
Dateofjoining                0  
LastWorkingDate             17488  
Joining Designation          0  
Grade                        0  
Total Business Value         0  
Quarterly Rating             0  
dtype: int64
```

Presence of null value. To be handled

```
In [12]: immuting_df=df.drop(columns=['MMM-YY', 'LastWorkingDate', 'Dateofjoining', 'City'])
```

```
In [13]: from sklearn.impute import KNNImputer

imputer = KNNImputer(n_neighbors=3)
```

```
In [14]: col=immuting_df.columns
immuting_df=imputer.fit_transform(immuting_df)
immuting_df=pd.DataFrame(immuting_df,columns=col)
```

```
In [15]: df['Age']=immuting_df['Age']
```

```
In [16]: df.isnull().sum()
```

```
Out[16]: MMM-YY                0
Driver_ID                    0
Age                          0
Gender                       52
City                         0
Education_Level              0
Income                      0
Dateofjoining                0
LastWorkingDate             17488
Joining Designation          0
Grade                       0
Total Business Value         0
Quarterly Rating             0
dtype: int64
```

Label Encoding of city

```
In [17]: from sklearn.preprocessing import LabelEncoder
```

```
In [18]: le = LabelEncoder()
df['City']=le.fit_transform(df['City'])
```

```
In [19]: df.head()
```

```
Out[19]:
```

	MMM-YY	Driver_ID	Age	Gender	City	Education_Level	Income	Dateofjoining	LastWorkingDate	Joining Designation	Grade	Total Business Value	Quarterly Rating
0	01/01/19	1	28.0	0.0	15	2	57387	24/12/18	NaN	1	1	2381060	2
1	02/01/19	1	28.0	0.0	15	2	57387	24/12/18	NaN	1	1	-665480	2
2	03/01/19	1	28.0	0.0	15	2	57387	24/12/18	03/11/19	1	1	0	2
3	11/01/20	2	31.0	0.0	26	2	67016	11/06/20	NaN	2	2	0	1
4	12/01/20	2	31.0	0.0	26	2	67016	11/06/20	NaN	2	2	0	1

Covertng to date time

```
In [20]: df['MMM-YY']=pd.to_datetime(df['MMM-YY'])
```

```
In [21]: df['Dateofjoining']=pd.to_datetime(df['Dateofjoining'])
```

```
In [22]: df['LastWorkingDate']=pd.to_datetime(df['LastWorkingDate'])
```

Aggregating at drivers level

```
In [23]: agg={
    'Age':['max'],
    'Gender':['first'],
    'City':['first'],
    'Education_Level':['last'],
    'Income':['mean'],
    'Dateofjoining':['min'],
    'LastWorkingDate':['max'],
    'Joining Designation':['min'],
    'Grade':['max'],
    'Total Business Value':['sum'],
    'Quarterly Rating':['mean']
}
```

```
In [24]: groups = df.groupby('Driver_ID')
```

```
In [25]: df_grp = df.groupby(by='Driver_ID', as_index=False).aggregate(agg)
df_grp.columns = df_grp.columns.droplevel(level=1)
```

```
In [26]: #df.drop(columns='Driver_ID', inplace=True)
```



```
In [27]: df_grp.head(10)
```

```
Out[27]:
```

	Driver_ID	Age	Gender	City	Education_Level	Income	Dateofjoining	LastWorkingDate	Joining Designation	Grade	Total Business Value	Quarterly Rating
0	1	28.0	0.0	15	2	57387.0	2018-12-24	2019-03-11	1	1	1715580	2.00000
1	2	31.0	0.0	26	2	67016.0	2020-11-06	NaT	2	2	0	1.00000
2	4	43.0	0.0	4	2	65603.0	2019-12-07	2020-04-27	2	2	350000	1.00000
3	5	29.0	0.0	28	0	46368.0	2019-01-09	2019-03-07	1	1	120360	1.00000
4	6	31.0	1.0	2	1	78728.0	2020-07-31	NaT	3	3	1265000	1.60000
5	8	34.0	0.0	11	0	70656.0	2020-09-19	2020-11-15	3	3	0	1.00000
6	11	28.0	1.0	10	2	42172.0	2020-12-07	NaT	1	1	0	1.00000
7	12	35.0	0.0	15	2	28116.0	2019-06-29	2019-12-21	1	1	2607180	2.50000
8	13	31.0	0.0	10	2	119227.0	2015-05-28	2020-11-25	1	4	10213040	1.26087
9	14	39.0	1.0	18	0	19734.0	2020-10-16	NaT	3	3	0	1.00000

Creating a column which tells whether the quarterly rating has increased for that driver - for those whose quarterly rating has increased we assign the value 1

```
In [28]: df_grp['Rating_Increase']=0
```

```
In [29]: groups = df.groupby('Driver_ID')
for name, group in groups:
    ratings = group['Quarterly Rating'].tolist()
    # check if the rating values are increasing or have at least one increase
    if any(x < y for x, y in zip(ratings, ratings[1:])):
        df_grp.loc[(df_grp['Driver_ID'] == name), 'Rating_Increase'] = 1
```

Creating a column which tells whether the quarterly rating has decreased for that driver - for those whose quarterly rating has decreased we assign the value 1

```
In [30]: df_grp['Rating_Decrease']=0
```

```
In [31]: groups = df.groupby('Driver_ID')
for name, group in groups:
    ratings = group['Quarterly Rating'].tolist()
    # check if the rating values are increasing or have at least one increase
    if any(x > y for x, y in zip(ratings, ratings[1:])):
        df_grp.loc[(df_grp['Driver_ID'] == name), 'Rating_Decrease'] = 1
```

Creating a column which tells whether the monthly income has increased for that driver - for those whose monthly income has increased we assign the value 1


```
In [32]: df_grp['Monthly_Income_Increase']=0
```

```
In [33]: groups = df.groupby('Driver_ID')
for name, group in groups:
    monthly_income = group['Income'].tolist()
    # check if the rating values are increasing or have at least one increase
    if any(x < y for x, y in zip(monthly_income, monthly_income[1:])):
        df_grp.loc[(df_grp['Driver_ID'] == name), 'Monthly_Income_Increase'] = 1
```

In [34]: df_grp.describe()

Out[34]:

	Driver_ID	Age	Gender	City	Education_Level	Income	Joining Designation	Grade	Total Business Value	Quarterl Ratin
count	2381.000000	2381.000000	2381.000000	2381.000000	2381.000000	2381.000000	2381.000000	2381.000000	2.381000e+03	2381.000000
mean	1397.559009	33.759765	0.410332	14.074339	1.00756	59232.460484	1.820244	2.097018	4.586742e+06	1.56630
std	806.161628	5.944614	0.491997	8.252167	0.81629	28298.214012	0.841433	0.941702	9.127115e+06	0.71965
min	1.000000	21.000000	0.000000	0.000000	0.00000	10747.000000	1.000000	1.000000	-1.385530e+06	1.00000
25%	695.000000	30.000000	0.000000	7.000000	0.00000	39104.000000	1.000000	1.000000	0.000000e+00	1.00000
50%	1400.000000	33.000000	0.000000	14.000000	1.00000	55285.000000	2.000000	2.000000	8.176800e+05	1.00000
75%	2100.000000	37.000000	1.000000	21.000000	2.00000	75835.000000	2.000000	3.000000	4.173650e+06	2.00000
max	2788.000000	58.000000	1.000000	28.000000	2.00000	188418.000000	5.000000	5.000000	9.533106e+07	4.00000



Duplicate value check

In [35]: df.shape

Out[35]: (19104, 13)

In [36]: df.drop_duplicates(inplace=True)

In [37]: df.shape

Out[37]: (19104, 13)

In []:

Adding feature of grade change

```
In [38]: df_grp['Grade_Change']=df_grp['Grade']-df_grp['Joining Designation']
```

Adding Target feature of churning

```
In [39]: df_grp['churn']=df_grp['LastWorkingDate'].isnull().astype(int)
# 1: not resigned
# 2: resigned
```

Adding feature of tenure

```
In [40]: df_grp['tenure']=(df_grp['LastWorkingDate']-df_grp['Dateofjoining']).dt.days
```

```
In [41]: df_grp.drop(columns='LastWorkingDate',inplace=True)
```

```
In [42]: df_grp.drop(columns="Dateofjoining",inplace=True)
```

```
In [43]: df_grp.head()
```

Out[43]:


	Driver_ID	Age	Gender	City	Education_Level	Income	Joining Designation	Grade	Total Business Value	Quarterly Rating	Rating_Increase	Rating_Decrease	Monthly_
0	1	28.0	0.0	15	2	57387.0	1	1	1715580	2.0	0	0	
1	2	31.0	0.0	26	2	67016.0	2	2	0	1.0	0	0	
2	4	43.0	0.0	4	2	65603.0	2	2	350000	1.0	0	0	
3	5	29.0	0.0	28	0	46368.0	1	1	120360	1.0	0	0	
4	6	31.0	1.0	2	1	78728.0	3	3	1265000	1.6	1	0	

```
In [44]: df_grp['tenure'] = df_grp['tenure'].fillna(0)
```

Univariate Analysis

```
In [45]: sns.histplot(df_grp['Age'])
```


```
Out[45]: <AxesSubplot:xlabel='Age', ylabel='Count'>
```



```
In [46]: sns.countplot(df_grp['Gender'])
```

C:\Users\gokul\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(


```
Out[46]: <AxesSubplot:xlabel='Gender', ylabel='count'>
```



```
In [47]: sns.countplot(df_grp['Rating_Increase'])
```

C:\Users\gokul\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

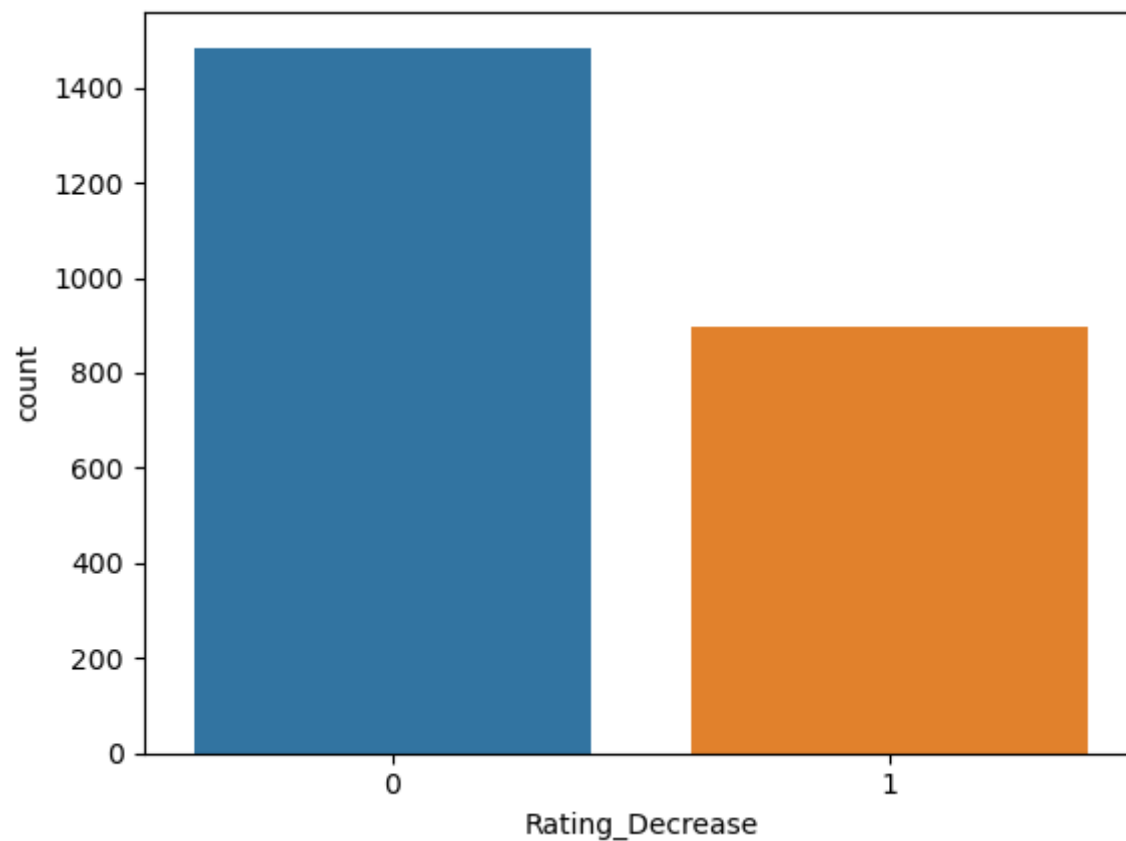
```
Out[47]: <AxesSubplot:xlabel='Rating_Increase', ylabel='count'>
```



```
In [48]: sns.countplot(df_grp['Rating_Decrease'])
```

C:\Users\gokul\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(


```
Out[48]: <AxesSubplot:xlabel='Rating_Decrease', ylabel='count'>
```




```
In [49]: sns.countplot(df_grp['Monthly_Income_Increase'])
```

C:\Users\gokul\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(


```
Out[49]: <AxesSubplot:xlabel='Monthly_Income_Increase', ylabel='count'>
```



```
In [50]: sns.countplot(df_grp['Education_Level'])
```


C:\Users\gokul\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

```
Out[50]: <AxesSubplot:xlabel='Education_Level', ylabel='count'>
```




```
In [51]: sns.histplot(df_grp['Income'])
```

```
Out[51]: <AxesSubplot:xlabel='Income', ylabel='Count'>
```



```
In [52]: sns.histplot(df_grp[df_grp['tenure']!=0]['tenure'])  
# during of the service
```

```
Out[52]: <AxesSubplot:xlabel='tenure', ylabel='Count'>
```

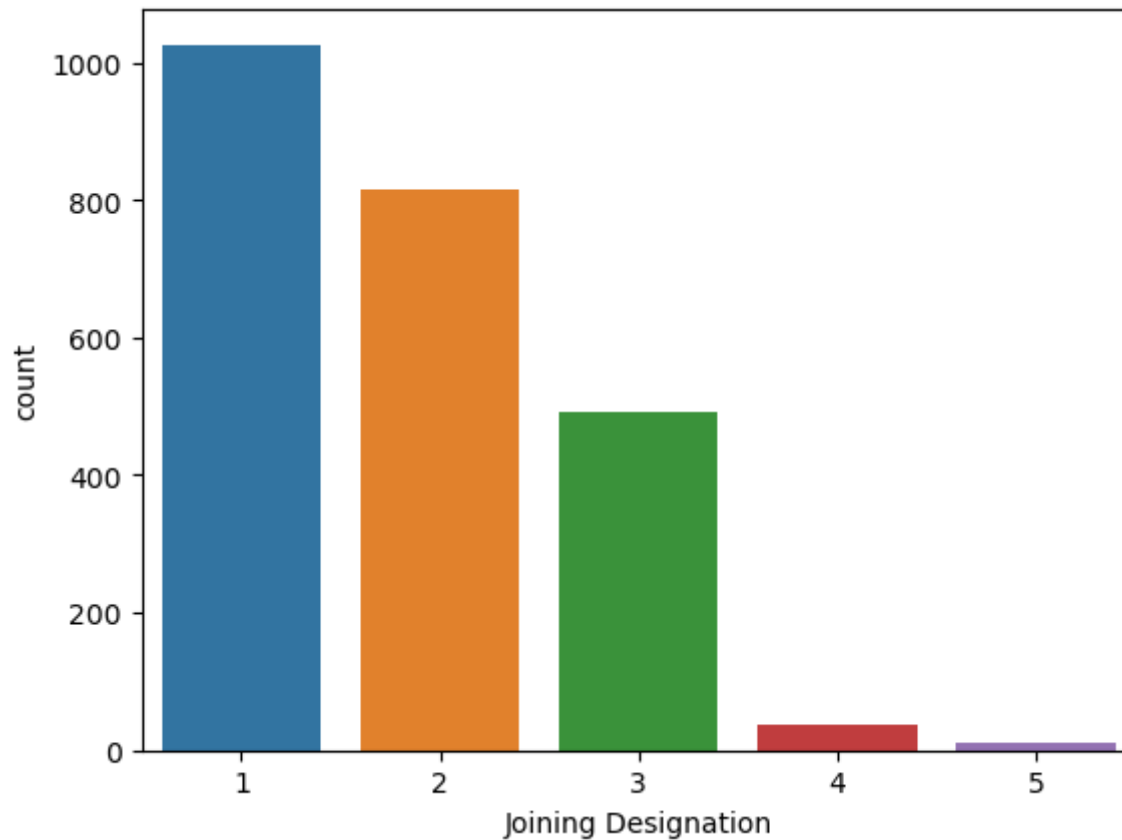


```
In [53]: sns.countplot(df_grp['Joining Designation'])
```

C:\Users\gokul\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```


```
Out[53]: <AxesSubplot:xlabel='Joining Designation', ylabel='count'>
```



```
In [54]: sns.countplot(df_grp['Grade'])
```

C:\Users\gokul\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(


```
Out[54]: <AxesSubplot:xlabel='Grade', ylabel='count'>
```



```
In [55]: sns.countplot(df_grp['churn'])
```


C:\Users\gokul\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(

```
Out[55]: <AxesSubplot:xlabel='churn', ylabel='count'>
```



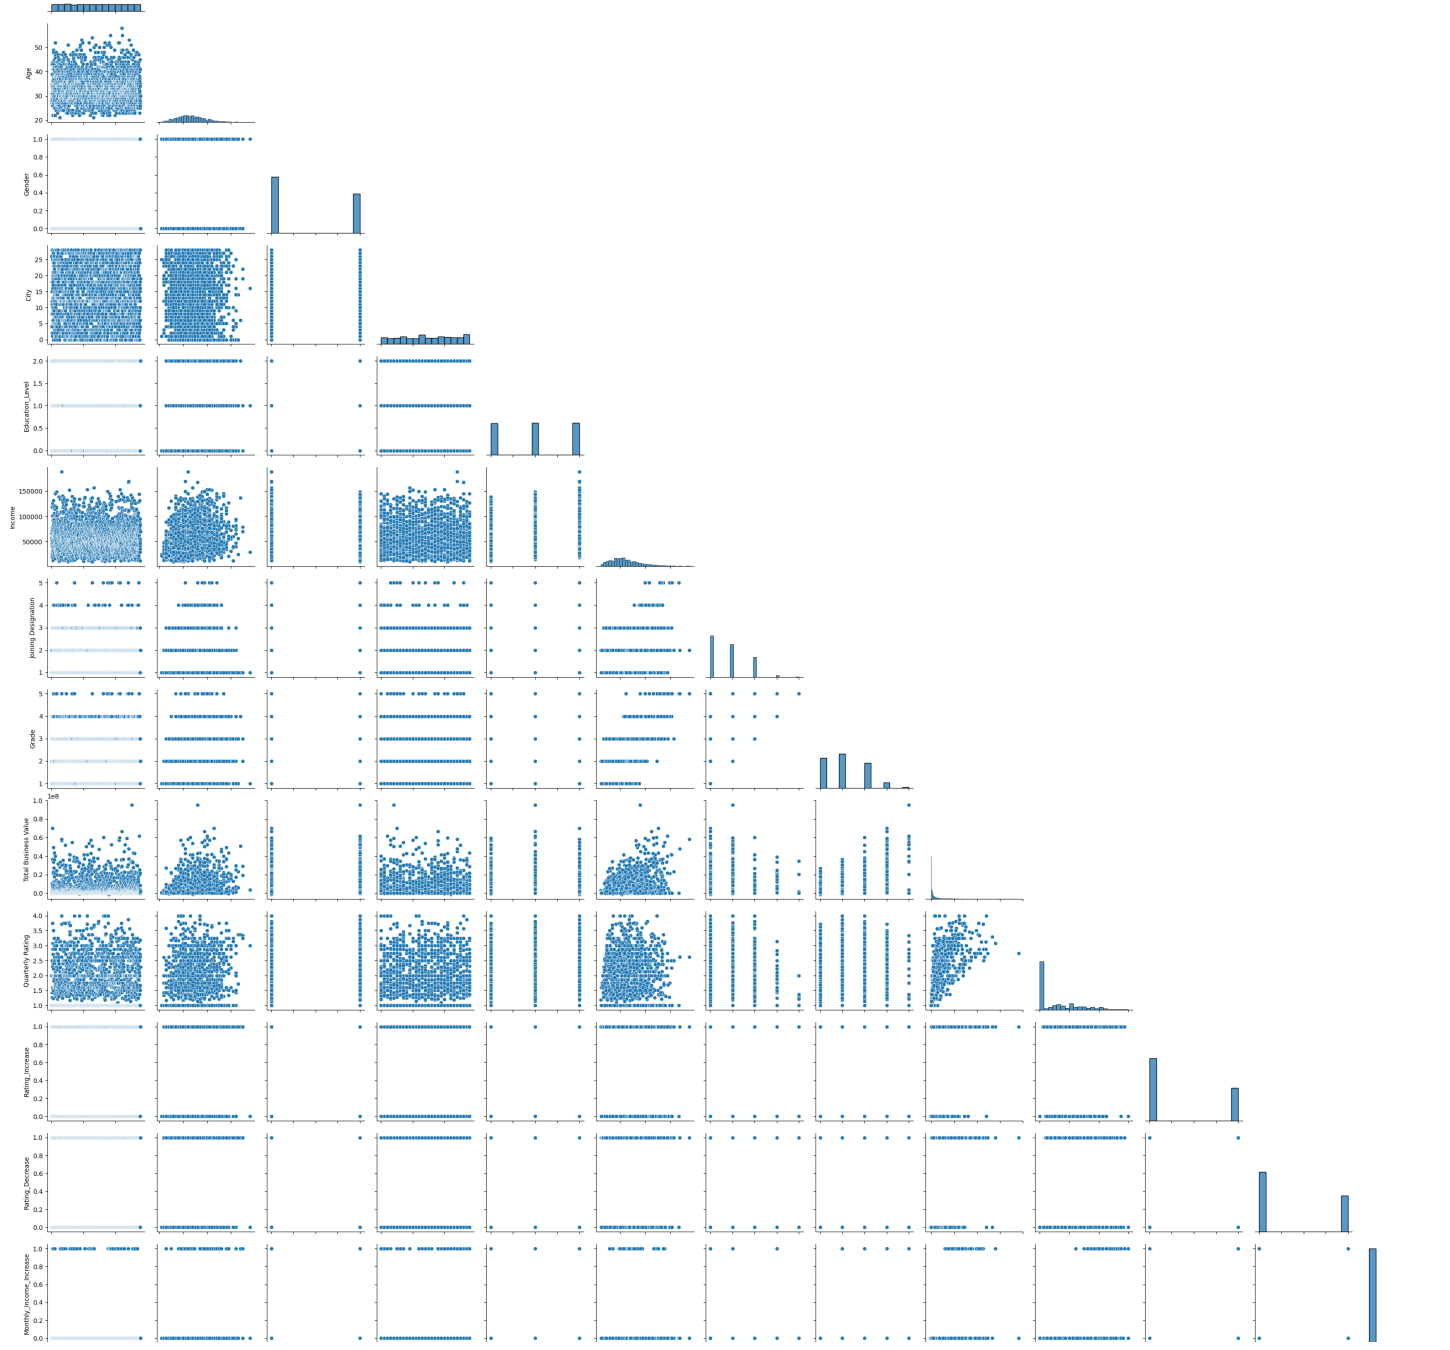
Bivariate Analysis


```
In [56]: plt.figure(figsize=(15,5))
sns.heatmap(df_grp.corr(),annot=True)
plt.show()
```




```
In [57]: sns.pairplot(df_grp, corner=True)
```

```
Out[57]: <seaborn.axisgrid.PairGrid at 0x1c525795d60>
```



- There are almost same number of people for high and low educated people are working as drivers.
- Income is left skewed. May need transformation
- People leaving the job, mostly left the before completing 400 days of work
- More people join at lower designation. Try giving higher designation when people with experience is joining
- Most people are in grade 2 designation
- Quarterly rating and Business value generate have high correlation
- Joining designation and Grade have high correlation. Which means people are stuck in their initial grade
- Income and grade have high correlation. So people are stuck in their first grade their income is also not rising
- If business values generated is high grade increased rapidly
- If business values generated is low grade can decrease

Model Building

Data Processing

In [58]: `df_grp.columns`

Out[58]: Index(['Driver_ID', 'Age', 'Gender', 'City', 'Education_Level', 'Income',
'Joining Designation', 'Grade', 'Total Business Value',
'Quarterly Rating', 'Rating_Increase', 'Rating_Decrease',
'Monthly_Income_Increase', 'Grade_Change', 'churn', 'tenure'],
dtype='object')

```
In [59]: df_grp.drop(columns=['tenure'],inplace=True)
```


```
In [60]: # Data Processing  
from sklearn.preprocessing import LabelEncoder  
from sklearn.preprocessing import StandardScaler  
from sklearn.model_selection import train_test_split  
from statsmodels.tools.tools import add_constant
```

Log tranfomation of income

```
In [61]: df_grp['Income']=np.log(df_grp['Income'])
```

```
In [62]: sns.histplot(df_grp['Income'])
```


```
Out[62]: <AxesSubplot:xlabel='Income', ylabel='Count'>
```



Outlier treatment

```
In [63]: sns.boxplot(data=df_grp, y='Income')
```

```
Out[63]: <AxesSubplot:ylabel='Income'>
```



```
In [64]: Q1 = np.percentile(df_grp['Income'], 25, interpolation = 'midpoint')
Q2 = np.percentile(df_grp['Income'], 50, interpolation = 'midpoint')
Q3 = np.percentile(df_grp['Income'], 75, interpolation = 'midpoint')
IQR = Q3 - Q1
low_lim = Q1 - 1.5 * IQR
up_lim = Q3 + 1.5 * IQR

df_grp=df_grp[(df_grp['Income']>low_lim) & (df_grp['Income']<up_lim)]
```

C:\Users\gokul\AppData\Local\Temp\ipykernel_4860\2903179881.py:1: DeprecationWarning: the `interpolation=` argument to percentile was renamed to `method=`, which has additional options.
Users of the modes 'nearest', 'lower', 'higher', or 'midpoint' are encouraged to review the method they. (Deprecated NumPy 1.22)

```
Q1 = np.percentile(df_grp['Income'], 25, interpolation = 'midpoint')
```

C:\Users\gokul\AppData\Local\Temp\ipykernel_4860\2903179881.py:2: DeprecationWarning: the `interpolation=` argument to percentile was renamed to `method=`, which has additional options.
Users of the modes 'nearest', 'lower', 'higher', or 'midpoint' are encouraged to review the method they. (Deprecated NumPy 1.22)

```
Q2 = np.percentile(df_grp['Income'], 50, interpolation = 'midpoint')
```

C:\Users\gokul\AppData\Local\Temp\ipykernel_4860\2903179881.py:3: DeprecationWarning: the `interpolation=` argument to percentile was renamed to `method=`, which has additional options.
Users of the modes 'nearest', 'lower', 'higher', or 'midpoint' are encouraged to review the method they. (Deprecated NumPy 1.22)

```
Q3 = np.percentile(df_grp['Income'], 75, interpolation = 'midpoint')
```


Independent and target feature split

```
In [65]: y=df_grp['churn']
X=df_grp.drop(['churn'], axis=1)
```


In [66]: X.head()

Out[66]:

	Driver_ID	Age	Gender	City	Education_Level	Income	Joining Designation	Grade	Total Business Value	Quarterly Rating	Rating_Increase	Rating_Decrease	Monthl
0	1	28.0	0.0	15	2	10.957573	1	1	1715580	2.0	0	0	
1	2	31.0	0.0	26	2	11.112687	2	2	0	1.0	0	0	
2	4	43.0	0.0	4	2	11.091377	2	2	350000	1.0	0	0	
3	5	29.0	0.0	28	0	10.744365	1	1	120360	1.0	0	0	
4	6	31.0	1.0	2	1	11.273754	3	3	1265000	1.6	1	0	



VIF Check

```
In [67]: from statsmodels.stats.outliers_influence import variance_inflation_factor

vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

```
C:\Users\gokul\anaconda3\lib\site-packages\statsmodels\stats\outliers_influence.py:193: RuntimeWarning: divide by zero encountered in double_scalars
  vif = 1. / (1. - r_squared_i)
```

Out[67]:

	Features	VIF
6	Joining Designation	inf
7	Grade	inf
13	Grade_Change	inf
5	Income	62.57
1	Age	36.78
9	Quarterly Rating	18.86
8	Total Business Value	4.50
0	Driver_ID	4.02
3	City	3.93
11	Rating_Decrease	3.03
10	Rating_Increase	2.95
4	Education_Level	2.65
2	Gender	1.71
12	Monthly_Income_Increase	1.28

```
In [68]: X.drop(columns=['Joining Designation','Income','Grade','Quarterly Rating'],inplace=True)
```

```
In [69]: vif = pd.DataFrame()
vif['Features'] = X.columns
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
vif['VIF'] = round(vif['VIF'], 2)
vif = vif.sort_values(by = "VIF", ascending = False)
vif
```

```
Out[69]:
```

	Features	VIF
1	Age	8.10
0	Driver_ID	3.73
3	City	3.67
5	Total Business Value	3.30
7	Rating_Decrease	2.59
6	Rating_Increase	2.55
4	Education_Level	2.49
9	Grade_Change	2.21
2	Gender	1.69
8	Monthly_Income_Increase	1.26

Test Train Split

```
In [70]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.6)
X_val, X_test, y_val, y_test = train_test_split(X_test, y_test, test_size=0.5)
```

Scaling - Using StandardScaler

```
In [71]: sc = StandardScaler()
sc.fit(X_train)
X_train = sc.transform(X_train)
X_test = sc.transform(X_test)
```

Handling Imbalance data

```
In [72]: from imblearn.over_sampling import SMOTE
sm = SMOTE()
X_train, y_train = sm.fit_resample(X_train, y_train)
X_val, y_val = sm.fit_resample(X_val, y_val)
```

Model Building-1 (Bagging Algorithm)

```
In [73]: from sklearn.ensemble import RandomForestClassifier
```

```
In [74]: model1=RandomForestClassifier()
```

```
In [75]: model1.fit(X_train, y_train)
```

```
Out[75]: RandomForestClassifier()
```

```
In [76]: y_hat_train=model1.predict(X_train)
y_hat_val=model1.predict(X_val)
```

```
C:\Users\gokul\anaconda3\lib\site-packages\sklearn\base.py:443: UserWarning: X has feature names, but RandomForestClassifier was fitted without feature names
  warnings.warn(
```

```
In [77]: train_score=f1_score(y_train,y_hat_train)
val_score=f1_score(y_val,y_hat_val)
print(train_score,val_score)
```

```
1.0 0.6736842105263158
```

Hyperparameter tuning

```
In [78]: from sklearn.model_selection import GridSearchCV
```

```
In [79]: model=RandomForestClassifier()
parameters = {'n_estimators':[50,100,150,200],
              'criterion':['gini', 'entropy', 'log_loss'],
              'max_depth':[5,6,7,8,9,10],
              'min_samples_split':[4,5,6,7]
              }
```

```
In [80]: clf = GridSearchCV(model, parameters, scoring = 'accuracy',cv = 3,n_jobs=-1)
```

```
In [81]: clf.fit(X_val, y_val)
```

```
C:\Users\gokul\anaconda3\lib\site-packages\sklearn\model_selection\_validation.py:372: FitFailedWarning:
288 fits failed out of a total of 864.
```

```
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

```
Below are more details about the failures:
```

```
-----
288 fits failed with the following error:
```

```
Traceback (most recent call last):
```

```
File "C:\Users\gokul\anaconda3\lib\site-packages\sklearn\model_selection\_validation.py", line 680, in _fit_and_score
```

```
    estimator.fit(X_train, y_train, **fit_params)
```

```
File "C:\Users\gokul\anaconda3\lib\site-packages\sklearn\ensemble\_forest.py", line 450, in fit
```

```
    trees = Parallel(
```

```
File "C:\Users\gokul\anaconda3\lib\site-packages\joblib\parallel.py", line 1085, in __call__
```

```
    if self.dispatch_one_batch(iterator):
```

```
File "C:\Users\gokul\anaconda3\lib\site-packages\joblib\parallel.py", line 901, in dispatch_one_batch
```

```
    self._dispatch(tasks)
```

```
File "C:\Users\gokul\anaconda3\lib\site-packages\joblib\parallel.py", line 819, in _dispatch
```

```
In [82]: clf.best_estimator_
```

```
Out[82]: RandomForestClassifier(max_depth=10, min_samples_split=4)
```

```
In [83]: model1_val=clf.best_estimator_
```

```
In [84]: model1_val.fit(X_train, y_train)
```

```
Out[84]: RandomForestClassifier(max_depth=10, min_samples_split=4)
```

```
In [85]: y_hat_train=model1_val.predict(X_train)
y_hat_test=model1_val.predict(X_test)
```

```
In [86]: train_score_f1=f1_score(y_train,y_hat_train)
test_score_f1=f1_score(y_test,y_hat_test)
print(train_score_f1,test_score_f1)
```

```
0.9238493723849373 0.6047619047619047
```

```
In [87]: train_score=accuracy_score(y_train,y_hat_train)
test_score=accuracy_score(y_test,y_hat_test)
print(train_score,test_score)
```

```
0.927547770700637 0.7645390070921986
```

```
In [88]: feature_improtance=pd.DataFrame(model1_val.feature_importances_,X.columns)
feature_improtance.sort_values(0)
```

```
Out[88]:
```

	0
Monthly_Income_Increase	0.013941
Gender	0.018851
Grade_Change	0.047569
Education_Level	0.060319
Rating_Decrease	0.061729
Age	0.118673
Driver_ID	0.125464
City	0.133173
Rating_Increase	0.145645
Total Business Value	0.274635

```
In [89]: fpr, tpr, thresholds = roc_curve(y_test, y_hat_test)
precision, recall, thresholds = precision_recall_curve(y_test, y_hat_test)
```


```
In [90]: print('fpr, tpr, thresholds', fpr, tpr, thresholds)
print('precision, recall, thresholds', precision, recall, thresholds)
```

```
fpr, tpr, thresholds [0.          0.13807531 1.          ] [0.          0.55947137 1.          ] [0 1]
precision, recall, thresholds [0.32198582 0.65803109 1.          ] [1.          0.55947137 0.          ] [0 1]
```

Precision Recall Curve

```
In [91]: sns.lineplot(y=precision,x=recall)
```


```
Out[91]: <AxesSubplot:>
```



ROC AUC Curve

In [92]: `sns.lineplot(x=fpr,y=tpr)`

Out[92]: `<AxesSubplot:>`



Classification Report

```
In [93]: target_names = ['class 0', 'class 1']  
print(classification_report(y_test, y_hat_test, target_names=target_names))
```

	precision	recall	f1-score	support
class 0	0.80	0.86	0.83	478
class 1	0.66	0.56	0.60	227
accuracy			0.76	705
macro avg	0.73	0.71	0.72	705
weighted avg	0.76	0.76	0.76	705

Confusion matrix

```
In [94]: tn, fp, fn, tp = confusion_matrix(y_test, y_hat_test).ravel()  
  
print((tn, fp, fn, tp))  
  
(412, 66, 100, 127)
```

Comments on Metrics

- Area of precision recall curve is slightly above 0.5
- False negative are really large for this model
- Area under AUC Curve is largely 0.5
- After 0.5 trp for every increase in trp, there is a larger increade in fpr
- Accuracy of this model is around 76%
- F1 score of this model is around 0.60 which is really poor

Model Building-2 (Boosting Algorithm)

```
In [95]: import xgboost as xgb
```

```
In [96]: model2 = xgb.XGBClassifier()  
model2.fit(X_train, y_train)
```

```
Out[96]: XGBClassifier(base_score=None, booster=None, callbacks=None,  
                        colsample_bylevel=None, colsample_bynode=None,  
                        colsample_bytree=None, early_stopping_rounds=None,  
                        enable_categorical=False, eval_metric=None, feature_types=None,  
                        gamma=None, gpu_id=None, grow_policy=None, importance_type=None,  
                        interaction_constraints=None, learning_rate=None, max_bin=None,  
                        max_cat_threshold=None, max_cat_to_onehot=None,  
                        max_delta_step=None, max_depth=None, max_leaves=None,  
                        min_child_weight=None, missing=nan, monotone_constraints=None,  
                        n_estimators=100, n_jobs=None, num_parallel_tree=None,  
                        predictor=None, random_state=None, ...)
```

```
In [97]: y_hat_train2=model1.predict(X_train)  
y_hat_test2=model1.predict(X_test)
```

```
In [98]: train_score=f1_score(y_train,y_hat_train2)  
test_score=f1_score(y_test,y_hat_test2)  
print(train_score,test_score)
```

```
1.0 0.5741626794258374
```

Hyperparameter tuning

```
In [99]: model=xgb.XGBClassifier()  
parameters = {'subsample':[0.5,0.6,0.7,0.8,0.9,0.1],  
              'lambda':[0.01, 0.1, 1],  
              'alpha':[0.01, 0.1, 1],  
              'eta':[0,0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1]}  
}
```

```
In [100]: clf = GridSearchCV(model, parameters, scoring = 'accuracy',cv = 3,n_jobs=-1)
```

```
In [101]: clf.fit(X_val, y_val)
```

```
Out[101]: GridSearchCV(cv=3,  
                       estimator=XGBClassifier(base_score=None, booster=None,  
                                                callbacks=None, colsample_bylevel=None,  
                                                colsample_bynode=None,  
                                                colsample_bytree=None,  
                                                early_stopping_rounds=None,  
                                                enable_categorical=False, eval_metric=None,  
                                                feature_types=None, gamma=None,  
                                                gpu_id=None, grow_policy=None,  
                                                importance_type=None,  
                                                interaction_constraints=None,  
                                                learning_rate=None, ...  
                                                max_delta_step=None, max_depth=None,  
                                                max_leaves=None, min_child_weight=None,  
                                                missing=nan, monotone_constraints=None,  
                                                n_estimators=100, n_jobs=None,  
                                                num_parallel_tree=None, predictor=None,  
                                                random_state=None, ...),  
                       n_jobs=-1,  
                       param_grid={'lambda': [0.01, 0.1, 1],  
                                   'alpha': [0.01, 0.1, 1],  
                                   'eta': [0, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1],  
                                   'subsample': [0.5, 0.6, 0.7, 0.8, 0.9, 0.1]})
```

```
In [102]: clf.best_estimator_
```

```
Out[102]: XGBClassifier(alpha=0.01, base_score=None, booster=None, callbacks=None,
                        colsample_bylevel=None, colsample_bynode=None,
                        colsample_bytree=None, early_stopping_rounds=None,
                        enable_categorical=False, eta=0.1, eval_metric=None,
                        feature_types=None, gamma=None, gpu_id=None, grow_policy=None,
                        importance_type=None, interaction_constraints=None, lambda=0.1,
                        learning_rate=None, max_bin=None, max_cat_threshold=None,
                        max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
                        max_leaves=None, min_child_weight=None, missing=nan,
                        monotone_constraints=None, n_estimators=100, n_jobs=None, ...)
```

```
In [103]: model2_val=clf.best_estimator_
```

```
In [104]: model2_val.fit(X_train, y_train)
```

```
Out[104]: XGBClassifier(alpha=0.01, base_score=None, booster=None, callbacks=None,
                        colsample_bylevel=None, colsample_bynode=None,
                        colsample_bytree=None, early_stopping_rounds=None,
                        enable_categorical=False, eta=0.1, eval_metric=None,
                        feature_types=None, gamma=None, gpu_id=None, grow_policy=None,
                        importance_type=None, interaction_constraints=None, lambda=0.1,
                        learning_rate=None, max_bin=None, max_cat_threshold=None,
                        max_cat_to_onehot=None, max_delta_step=None, max_depth=None,
                        max_leaves=None, min_child_weight=None, missing=nan,
                        monotone_constraints=None, n_estimators=100, n_jobs=None, ...)
```

```
In [105]: feature_improtance=pd.DataFrame(model2_val.feature_importances_,X.columns)
feature_improtance.sort_values(0)
```

```
Out[105]:
```

	0
Monthly_Income_Increase	0.012569
Driver_ID	0.045479
Gender	0.052563
City	0.054053
Age	0.054864
Education_Level	0.065184
Total Business Value	0.082039
Grade_Change	0.114659
Rating_Decrease	0.172166
Rating_Increase	0.346423

```
In [106]: y_hat_train=model2_val.predict(X_train)
y_hat_test=model2_val.predict(X_test)
```

```
In [107]: train_score_f1=f1_score(y_train,y_hat_train)
test_score_f1=f1_score(y_test,y_hat_test)
print(train_score_f1,test_score_f1)
```

```
0.9830234438156832 0.580046403712297
```

```
In [108]: train_score=accuracy_score(y_train,y_hat_train)
test_score=accuracy_score(y_test,y_hat_test)
print(train_score,test_score)
```

```
0.98328025477707 0.7432624113475177
```

```
In [109]: fpr, tpr, thresholds = roc_curve(y_test, y_hat_test)
precision, recall, thresholds = precision_recall_curve(y_test, y_hat_test)
```


```
In [110]: print('fpr, tpr, thresholds', fpr, tpr, thresholds)
print('precision, recall, thresholds', precision, recall, thresholds)
```

```
fpr, tpr, thresholds [0.          0.16527197 1.          ] [0.          0.55066079 1.          ] [0 1]
precision, recall, thresholds [0.32198582 0.6127451 1.          ] [1.          0.55066079 0.          ] [0 1]
```

Precision Recall Curve

```
In [111]: sns.lineplot(y=precision,x=recall)
```


```
Out[111]: <AxesSubplot:>
```



ROC AUC Curve

In [112]: `sns.lineplot(x=fpr,y=tpr)`

Out[112]: `<AxesSubplot:>`



Classification Report

```
In [113]: target_names = ['class 0', 'class 1']  
print(classification_report(y_test, y_hat_test, target_names=target_names))
```

	precision	recall	f1-score	support
class 0	0.80	0.83	0.82	478
class 1	0.61	0.55	0.58	227
accuracy			0.74	705
macro avg	0.70	0.69	0.70	705
weighted avg	0.74	0.74	0.74	705

Confusion matrix

```
In [114]: tn, fp, fn, tp = confusion_matrix(y_test, y_hat_test).ravel()  
  
print((tn, fp, fn, tp))  
  
(399, 79, 102, 125)
```

Comments on Metrics

- Area of precision recall curve is 0.5
- False negative are really large for this model
- Area under AUC Curve is more than 0.5
- After 0.5 trp for every increase in trp, there is a larger increade in fpr
- Accuracy of this model is around 75%
- F1 score of this model is around 0.58 which is really poor

Actionable Insights

- Joining designation and Grade have high correlation. Which means people are stuck in their initial grade
- Income and grade have high correlation. So people are stuck in their first grade their income is also not rising
- Grade_Change, Rating_Decrease, Rating_Increase, Total Business Value are the factors affecting the churn
- Accuracy of Bagging algorithm is coming as 80% whereas for boosting it is 68%
- After 500 days with the company churn rate decreases drastically
- City, age and gender contribute to less churn

Recommendations

- Grade promotion criteria to be relaxed
- Rating_Decrease criteria to be made stringent
- Steps to be taken to improve the
- Giving more incentives to the drivers during the initial stage with the company. Try to make them to stay for more than 650 days.
- Use random forest for prediction because it is computationally less expensive
- Use the random forest model to predict whether the person will churn or not and give incentives of grade change for those drivers
- Burden on driver for accruing business to be reduced.
- Training to improve the performance of the driver done in order

In []: